

# Andrew Vu - CS156\_HW2

March 4, 2022

## 1 CS156 (Introduction to AI), Spring 2022

## 2 Homework 2 submission

**2.0.1 Roster Name:** Andrew Vu

**2.0.2 Student ID:** 015055911

**2.0.3 Email address:** andrew.k.vu@sjsu.edu

Any special notes or anything you would like to communicate to me about this homework submission goes in here.

### 2.1 References and sources

List all your references and sources here. This includes all sites/discussion boards/blogs/posts/etc. where you grabbed some code examples.

From files section of Canvas page:

knn.synthetic\_data

kmeans.synthetic\_data

<https://machinelearningmastery.com/tutorial-to-implement-k-nearest-neighbors-in-python-from-scratch/>

<https://towardsdatascience.com/how-to-build-knn-from-scratch-in-python-5e22b8920bd2>

<https://www.askpython.com/python/examples/k-nearest-neighbors-from-scratch>

<https://medium.com/analytics-vidhya/knn-implementation-from-scratch-96-6-accuracy-python-machine-learning-31ba66958644>

<https://medium.com/analytics-vidhya/implementing-k-nearest-neighbours-knn-without-using-scikit-learn-3905b4decc3c>

### 2.2 Solution

Load libraries and set random number generator seed

```
[237]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from scipy.spatial import distance
from scipy.stats import mode
from sklearn.metrics import accuracy_score
```

```
[238]: np.random.seed(42)
```

Code the solution

### 2.2.1 KNN FUNCTION

```
[239]: # newObservation is a point in the test data set
# referenceData is the training data set

def knn(newObservation, referenceData, k=3):
    distances_from_newdata = [] # array to store distances
    k_shortest_distances = [] # array for the k shortest distances

    # compute the distance between each point in the reference data and the
    ↪ newObservation point
    for i in range(len(referenceData)):
        dist = distance.euclidean(referenceData.iloc[i], newObservation)
        distances_from_newdata.append([dist, i]) # add this computed distance
    ↪ to the distances_from_newdata array with its index
    distances_from_newdata.sort()
    k_shortest_distances = (distances_from_newdata)[0:k]

    return k_shortest_distances # returns the array of the k shortest distances
    ↪ with their respective row number
```

### 2.2.2 PREDICTIONS FUNCTION

```
[240]: def predictions(X_test, ref_data, Y_train):
    labels = [] # output
    for i in range(len(X_test)):
        modelabels = [] # create a separate array for determining the majority
    ↪ of the labels
        distance_indexes = knn(X_test.iloc[i], ref_data, k=3)
        for j in range(len(distance_indexes)): # for however many distances
    ↪ there are in the return value of knn
            modelabels.append(Y_train[distance_indexes[j][1]]) # append that
    ↪ distance's Y_train label value to a separate array
        lab = mode(modelabels) # returns an array of the majority of the labels
    ↪ as well as the count of how many
        lab = lab.mode[0] # returns specifically the mode
        labels.append(lab) # adds this label to the real labels
    return labels
```

## 2.3 Generate 2D Data

```
[241]: n = 100 # n = data points

# loc = mean or mass of gravity in which the generated data centers around this_
# value
# scale = standard deviation
# size = how many data points you want

# CLASS 0
X0 = np.random.normal(loc = -2.0, scale = 2.0, size = int(n/2))
Y0 = np.random.normal(loc = 0.0, scale = 1.0, size = int(n/2))

# CLASS 1
X1 = np.random.normal(loc = 2.0, scale = 2.0, size = int(n/2))
Y1 = np.random.normal(loc = 0.0, scale = 1.0, size = int(n/2))

X = np.concatenate((X0, X1), axis = 0)
Y = np.concatenate((Y0, Y1), axis = 0)
```

### 2.3.1 Generate Labels for 2D

```
[242]: l1 = [0]*int(n/2)
l2 = [1]*int(n/2)
labels2D = l1+l2
```

### 2.3.2 Visualize the generated data

```
[243]: dt = pd.DataFrame({'X':X, 'Y':Y}, columns=['X', 'Y'])
dt.head()
```

```
[243]:
```

	X	Y
0	-1.006572	0.324084
1	-2.276529	-0.385082
2	-0.704623	-0.676922
3	1.046060	0.611676
4	-2.468307	1.031000

```
[244]: X_train, X_test, Y_train, Y_test = train_test_split(dt, labels2D, test_size=0.
# 2, random_state=0)
# X_train = data for training set
# X_test = data for test set
# Y_train = labels 0/1 for training set
# Y_test = labels 0/1 or test set
# 80% training, 20% test split
print(X_train)
```

	X	Y
43	-2.602207	-0.327662
62	2.120460	1.158596
3	1.046060	0.611676
71	0.181225	-0.815810
45	-3.439688	-1.463515
..	...	...
96	-0.473901	-0.883857
67	-0.337356	1.896793
64	1.615278	0.963376
47	0.114244	0.261055
44	-4.957044	-0.392108

[80 rows x 2 columns]

### 2.3.3 ACCURACY FOR 2D

```
[245]: prediction_results = predictions(X_test, X_train, Y_train)
print(Y_test)
print(prediction_results)
print("Accuracy of the predictions on the 2D test dataset is " +
      ↪str(accuracy_score(Y_test, prediction_results)))
```

[0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0]

[0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0]


Accuracy of the predictions on the 2D test dataset is 0.9

### 2.3.4 ACTUAL VS PREDICTED LABELS VISUALIZATION

```
[246]: plt.subplot(1, 2, 1)
plt.scatter(X_train.iloc[:,0],X_train.iloc[:,1], s=25, c=Y_train, marker=".")
plt.scatter(X_test.iloc[:,0],X_test.iloc[:,1], s=50, c=Y_test, marker="v")
plt.title("Actual labels")

plt.subplot(1, 2, 2)
plt.scatter(X_train.iloc[:,0],X_train.iloc[:,1], s=25, c=Y_train, marker=".")
plt.scatter(X_test.iloc[:,0],X_test.iloc[:,1], s=50, c=prediction_results,
      ↪marker="v")
plt.title("Predicted labels")

plt.tight_layout()
plt.show()
```



## 2.4 Generate 3-D data

```
[247]: n = 1000 # n data points

# CLASS 0
X0 = np.random.normal(loc = 0, scale = 3, size = int(n/4))
Y0 = np.random.normal(loc = -3, scale = 1, size = int(n/4))
Z0 = np.random.normal(loc = -1, scale = 1, size = int(n/4))
# CLASS 1
X1 = np.random.normal(loc = 0, scale = 3, size = int(n/4))
Y1 = np.random.normal(loc = 1, scale = 2, size = int(n/4))
Z1 = np.random.normal(loc = 1, scale = 1, size = int(n/4))
# CLASS 2
X2 = np.random.normal(loc = 0, scale = 3, size = int(n/4))
Y2 = np.random.normal(loc = 3, scale = 1, size = int(n/4))
Z2 = np.random.normal(loc = 4, scale = 1, size = int(n/4))
# CLASS 3
X3 = np.random.normal(loc = 0, scale = 3, size = int(n/4))
Y3 = np.random.normal(loc = 5, scale = 3, size = int(n/4))
Z3 = np.random.normal(loc = -3, scale = 1, size = int(n/4))

X = np.concatenate((X0, X1, X2, X3), axis = 0)
Y = np.concatenate((Y0, Y1, Y2, Y3), axis = 0)
```

```
Z = np.concatenate((Z0, Z1, Z2, Z3), axis = 0)
```

```
[248]: l1 = [0]*int(n/4)
l2 = [1]*int(n/4)
l3 = [2]*int(n/4)
l4 = [3]*int(n/4)
labels3D = l1+l2+l3+l4
```

```
[249]: dt = pd.DataFrame({'X':X, 'Y':Y, 'Z':Z}, columns=['X', 'Y', 'Z'])
dt.head()
```

```
[249]:
```

	X	Y	Z
0	1.073362	-3.062679	-1.522723
1	1.682354	-2.044858	0.049009
2	3.249154	-3.985726	-1.704344
3	3.161406	-2.495953	-2.408461
4	-4.133008	-3.530258	-2.556629

```
[250]: X_train, X_test, Y_train, Y_test = train_test_split(dt, labels3D, test_size=0.
↳2, random_state=0)
print(X_train)
```

	X	Y	Z
687	1.272183	4.549020	3.075767
500	1.051890	1.226968	5.804348
332	1.197669	2.416217	1.323168
979	-0.790345	7.956349	-2.484372
817	-6.657901	6.364223	-5.832156
..	...	...	...
835	-0.885270	3.944236	-2.553127
192	-3.960700	-3.259591	0.066675
629	0.056549	2.549811	4.849102
559	3.944743	3.937570	4.198948
684	-0.079217	2.645959	4.032797

[800 rows x 3 columns]

## 2.5 ACCURACY FOR 3D

```
[251]: prediction_results = predictions(X_test, X_train, Y_train)
print(Y_test)
print("\n")
print(prediction_results)
print("Accuracy of the predictions on the 3D test dataset is " +
↳str(accuracy_score(Y_test, prediction_results)))
```

```
[3, 3, 1, 2, 2, 3, 0, 0, 1, 2, 1, 2, 3, 0, 2, 0, 3, 0, 0, 0, 0, 1, 1, 1, 3, 3,
0, 0, 3, 0, 3, 2, 2, 3, 1, 2, 3, 1, 3, 1, 0, 3, 2, 2, 1, 3, 3, 0, 1, 2, 1, 1, 1,
```

```
1, 3, 2, 1, 1, 3, 1, 2, 2, 0, 0, 1, 0, 1, 2, 1, 2, 2, 1, 2, 0, 3, 1, 1, 3, 2, 2,
1, 1, 1, 3, 3, 2, 0, 2, 1, 3, 1, 1, 3, 1, 2, 1, 2, 0, 1, 2, 3, 3, 0, 3, 1, 2, 2,
3, 3, 0, 0, 2, 1, 2, 1, 1, 2, 1, 1, 2, 3, 3, 0, 3, 0, 3, 1, 1, 3, 2, 1, 2, 2, 2,
3, 2, 1, 0, 2, 3, 0, 2, 2, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 3, 3, 3, 1, 1, 3, 1, 2,
3, 2, 3, 3, 1, 1, 3, 1, 1, 3, 3, 0, 3, 1, 3, 0, 2, 3, 0, 3, 3, 1, 1, 3, 3, 2, 1,
3, 0, 3, 3, 1, 0, 3, 2, 2, 0, 2, 0]
```

```
[3, 3, 1, 2, 2, 3, 0, 0, 2, 2, 1, 1, 3, 0, 2, 1, 3, 0, 0, 0, 0, 1, 1, 1, 3, 3,
0, 0, 3, 0, 3, 2, 2, 3, 0, 2, 3, 1, 3, 1, 0, 3, 2, 2, 1, 3, 3, 0, 1, 2, 1, 1, 1,
1, 3, 2, 1, 1, 3, 2, 2, 2, 0, 0, 1, 0, 1, 2, 1, 2, 2, 1, 2, 0, 3, 1, 1, 3, 2, 2,
1, 0, 1, 3, 3, 2, 0, 2, 1, 3, 1, 1, 3, 1, 2, 2, 2, 0, 1, 1, 3, 1, 0, 3, 1, 2, 2,
3, 3, 0, 0, 2, 1, 2, 1, 1, 2, 1, 1, 2, 3, 3, 0, 3, 0, 3, 1, 1, 3, 2, 1, 2, 2, 2,
3, 2, 1, 0, 2, 3, 0, 2, 2, 0, 1, 2, 1, 0, 0, 0, 1, 0, 1, 3, 3, 3, 1, 3, 3, 1, 1,
3, 2, 3, 3, 0, 2, 3, 1, 1, 3, 3, 0, 3, 1, 3, 0, 2, 3, 0, 3, 0, 1, 1, 3, 3, 2, 1,
3, 1, 3, 3, 1, 0, 3, 2, 2, 0, 2, 0]
```

Accuracy of the predictions on the 3D test dataset is 0.915

### 2.5.1 ACTUAL VS PREDICTED LABELS VISUALIZATION

```
[252]: plt.subplot(1, 2, 1)
plt.scatter(X_train.iloc[:,0],X_train.iloc[:,1], s=25, c=Y_train, marker=".")
plt.scatter(X_test.iloc[:,0],X_test.iloc[:,1], s=50, c=Y_test, marker="v")
plt.title("Actual labels")

plt.subplot(1, 2, 2)
plt.scatter(X_train.iloc[:,0],X_train.iloc[:,1], s=25, c=Y_train, marker=".")
plt.scatter(X_test.iloc[:,0],X_test.iloc[:,1], s=50, c=prediction_results,↵
↵marker="v")
plt.title("Predicted labels")

plt.tight_layout()
plt.show()
```

